

The University of Hong Kong
School of Computing and Data Science

Semantic-Motion

A Multi-view Engineering Pipeline for Perception, Fact-preserving Language
Augmentation and Fail-closed Multimodal Refinement

Author:	Your Name
Student ID:	Your Student ID
Supervisor:	Your Supervisor
Programme:	COMP7705 Project Report
Submission date:	Month 2026

Declaration

State originality, academic integrity, use of generative AI where required, and individual contributions for a group project.

Abstract

Keep the abstract below 200 words. State the data problem, the four-module pipeline, the evaluated outputs, the strongest quantitative result, and the main limitation. Do not describe planned work as completed work.

Contents

Declaration	i
Abstract	ii
1 Introduction	1
1.1 Project Context: From Motion Data to VLA-ready Multimodal Data	1
1.2 Engineering Problem and Motivation	1
1.3 Project Aim and Engineering Deliverables	1
1.4 System Requirements	1
1.5 Scope and Boundaries	1
1.6 Contributions	1
1.7 Report Organisation	1
2 Technical Background and Related Systems	2
2.1 Vision-Language-Action Models	2
2.2 Robot Learning Benchmarks	2
2.3 Vision-Language Models for Robotics	2
2.4 Video-to-Task Generation	2
2.5 Demonstration Data Quality	2
2.6 Engineering Gap	2
3 System Requirements and Overall Design	3
3.1 Inputs, Outputs and Use Cases	3
3.2 Design Goals	3
3.3 Overall Architecture	3
3.4 Module Responsibilities and Boundaries	3
3.5 Data Models and Interface Contracts	3
3.5.1 Paired Multi-view Video Contract	3
3.5.2 Shared Timebase and Motion Trajectory Contract	3
3.6 Observability and Visualisation Design	3
3.7 End-to-End Workflow	3
4 Design and Implementation of the Four Modules	4
4.1 The Rendering Module	4
4.1.1 Motion Data Import and Scene Initialisation	4
4.1.2 Scene Context Generation	4
4.1.3 Multi-view Video Synthesis	4

4.1.4	3D Motion Trajectory Export	4
4.1.5	Rendering Outputs and Manifest	4
4.2	The Perception Module	4
4.2.1	Paired-view Input and Synchronous Frame Sampling	4
4.2.2	Overlapping Macro Windows and Boundary Voting	4
4.2.3	VLM-based Structured Scene Understanding	4
4.2.4	Macro-Intent Recognition	4
4.2.5	Dense Micro-Instruction Parsing	4
4.2.6	Body, Contact, Posture and Trajectory Evidence	4
4.2.7	Temporal Task Boundary Detection	4
4.2.8	Segment-level Temporal Aggregation	4
4.2.9	Static Perception Benchmark	4
4.3	The Fact-preserving LLM Augmentation Module	5
4.3.1	Input and Interface	5
4.3.2	Instruction Reconstruction	5
4.3.3	Augmentation Strategies	5
4.3.4	Deterministic Fact Preservation and Hallucination Control	5
4.3.5	Output and Provenance	5
4.4	The Refinement Module	5
4.4.1	Sample Construction	5
4.4.2	Motion Data Normalisation	5
4.4.3	Motion Quality Scoring	5
4.4.4	Semantic Consistency Verification	5
4.4.5	Deterministic Cross-modal Alignment Checks	5
4.4.6	Fail-closed Refinement Decision and Explainable Outputs	5
5	System Integration, Demonstration and Evaluation	6
5.1	Research Questions	6
5.2	Integration Environment and Technology Stack	6
5.3	Module Integration and Interface Verification	6
5.4	End-to-End Demonstration and Visualisation	6
5.5	Evaluation Setup	6
5.6	Perception Module Evaluation	6
5.7	Pinned Video2Tasks Full-pipeline Comparison	6
5.8	Refinement and Alignment Evaluation	6
5.8.1	Human Gold Annotation Protocol	6
5.8.2	Temporal Boundary and Segment Metrics	6
5.8.3	Keep/Drop and Confidence Calibration	6
5.8.4	Controlled Motion Corruptions	6
5.9	Cross-module Case Studies	6

5.10	Component and Configuration Analysis	6
5.11	Engineering Acceptance Summary	6
6	Engineering Discussion	7
6.1	Interpretation of System Results	7
6.2	Engineering Strengths	7
6.3	Failure Mode Analysis	7
6.4	Limitations	7
6.5	Engineering Trade-offs	7
6.6	Evaluation Validity	7
6.7	Implications for VLA Data Pipelines	7
7	Conclusion and Future Work	8
7.1	Conclusion	8
7.2	Future Work	8
A	VLM System Prompts	9
B	Semantic Consistency Prompt	10
C	Data Schemas	11
D	Configuration	12
E	Additional Results	13
F	User Guide	14

List of Figures

List of Tables

Introduction

1.1 Project Context: From Motion Data to VLA-ready Multimodal Data

1.2 Engineering Problem and Motivation

1.3 Project Aim and Engineering Deliverables

1.4 System Requirements

1.5 Scope and Boundaries

1.6 Contributions

1.7 Report Organisation

Technical Background and Related Systems

2.1 Vision-Language-Action Models

2.2 Robot Learning Benchmarks

2.3 Vision-Language Models for Robotics

2.4 Video-to-Task Generation

2.5 Demonstration Data Quality

2.6 Engineering Gap

System Requirements and Overall Design

3.1 Inputs, Outputs and Use Cases

3.2 Design Goals

3.3 Overall Architecture

3.4 Module Responsibilities and Boundaries

3.5 Data Models and Interface Contracts

3.5.1 Paired Multi-view Video Contract

3.5.2 Shared Timebase and Motion Trajectory Contract

3.6 Observability and Visualisation Design

3.7 End-to-End Workflow

Design and Implementation of the Four Modules

4.1 The Rendering Module

4.1.1 Motion Data Import and Scene Initialisation

4.1.2 Scene Context Generation

4.1.3 Multi-view Video Synthesis

4.1.4 3D Motion Trajectory Export

4.1.5 Rendering Outputs and Manifest

4.2 The Perception Module

4.2.1 Paired-view Input and Synchronous Frame Sampling

4.2.2 Overlapping Macro Windows and Boundary Voting

4.2.3 VLM-based Structured Scene Understanding

4.2.4 Macro-Intent Recognition

4.2.5 Dense Micro-Instruction Parsing

4.2.6 Body, Contact, Posture and Trajectory Evidence

4.2.7 Temporal Task Boundary Detection

4.2.8 Segment-level Temporal Aggregation

4.2.9 Static Perception Benchmark

4.3 The Fact-preserving LLM Augmentation Module

4.3.1 Input and Interface

4.3.2 Instruction Reconstruction

4.3.3 Augmentation Strategies

4.3.4 Deterministic Fact Preservation and Hallucination Control

4.3.5 Output and Provenance

4.4 The Refinement Module

4.4.1 Sample Construction

4.4.2 Motion Data Normalisation

4.4.3 Motion Quality Scoring

4.4.4 Semantic Consistency Verification

4.4.5 Deterministic Cross-modal Alignment Checks

4.4.6 Fail-closed Refinement Decision and Explainable Outputs

System Integration, Demonstration and Evaluation

5.1 Research Questions

5.2 Integration Environment and Technology Stack

5.3 Module Integration and Interface Verification

5.4 End-to-End Demonstration and Visualisation

5.5 Evaluation Setup

5.6 Perception Module Evaluation

5.7 Pinned Video2Tasks Full-pipeline Comparison

5.8 Refinement and Alignment Evaluation

5.8.1 Human Gold Annotation Protocol

5.8.2 Temporal Boundary and Segment Metrics

5.8.3 Keep/Drop and Confidence Calibration

5.8.4 Controlled Motion Corruptions

5.9 Cross-module Case Studies

5.10 Component and Configuration Analysis

5.11 Engineering Acceptance Summary

Engineering Discussion

6.1 Interpretation of System Results

6.2 Engineering Strengths

6.3 Failure Mode Analysis

6.4 Limitations

6.5 Engineering Trade-offs

6.6 Evaluation Validity

6.7 Implications for VLA Data Pipelines

Conclusion and Future Work

7.1 Conclusion

Report only automatically verified or independently evaluated results. Mark API, Blender, and pending-human evidence explicitly; do not infer completion.

7.2 Future Work

VLM System Prompts

Semantic Consistency Prompt

Data Schemas

Configuration

Additional Results

User Guide
